

Heart_Attack_Project_Notebook

December 7, 2023

1 Imports

```
[1]: # !pip install catboost
     # !pip install xgboost
```

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Models
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from xgboost import XGBClassifier
from catboost import CatBoostClassifier

# Misc
from sklearn.metrics import accuracy_score, roc_auc_score, f1_score
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
import warnings

warnings.filterwarnings("ignore")
pd.set_option('display.max_rows', 500)
pd.set_option('display.max_columns', 500)
```

2 Introduction to the Data

```
[3]: df = pd.read_csv("heart_attack.csv")
```

```
[4]: df.shape
```

```
[4]: (8763, 26)
```

```
[5]: df.head()
```

```
[5]: Patient ID Age Sex Cholesterol Blood Pressure Heart Rate Diabetes \
0 BMW7812 67 Male 208 158/88 72 0
1 CZE1114 21 Male 389 165/93 98 1
2 BNI9906 21 Female 324 174/99 72 1
3 JLN3497 84 Male 383 163/100 73 1
4 GF08847 66 Male 318 91/88 93 1
```

```
Family History Smoking Obesity Alcohol Consumption \
0 0 1 0 0
1 1 1 1 1
2 0 0 0 0
3 1 1 0 1
4 1 1 1 0
```

```
Exercise Hours Per Week Diet Previous Heart Problems \
0 4.168189 Average 0
1 1.813242 Unhealthy 1
2 2.078353 Healthy 1
3 9.828130 Average 1
4 5.804299 Unhealthy 1
```

```
Medication Use Stress Level Sedentary Hours Per Day Income BMI \
0 0 9 6.615001 261404 31.251233
1 0 1 4.963459 285768 27.194973
2 1 9 9.463426 235282 28.176571
3 0 9 7.648981 125640 36.464704
4 0 6 1.514821 160555 21.809144
```

```
Triglycerides Physical Activity Days Per Week Sleep Hours Per Day \
0 286 0 6
1 235 1 7
2 587 4 4
3 378 3 4
4 231 1 5
```

```
Country Continent Hemisphere Heart Attack Risk
0 Argentina South America Southern Hemisphere 0
1 Canada North America Northern Hemisphere 0
2 France Europe Northern Hemisphere 0
3 Canada North America Northern Hemisphere 0
4 Thailand Asia Northern Hemisphere 0
```

```
[6]: df.describe()
```

```

[6]:
      Age Cholesterol Heart Rate Diabetes Family History \
count 8763.000000 8763.000000 8763.000000 8763.000000 8763.000000
mean  53.707977 259.877211 75.021682 0.652288 0.492982
std   21.249509 80.863276 20.550948 0.476271 0.499979
min   18.000000 120.000000 40.000000 0.000000 0.000000
25%   35.000000 192.000000 57.000000 0.000000 0.000000
50%   54.000000 259.000000 75.000000 1.000000 0.000000
75%   72.000000 330.000000 93.000000 1.000000 1.000000
max   90.000000 400.000000 110.000000 1.000000 1.000000

      Smoking Obesity Alcohol Consumption Exercise Hours Per Week \
count 8763.000000 8763.000000 8763.000000 8763.000000
mean  0.896839 0.501426 0.598083 10.014284
std   0.304186 0.500026 0.490313 5.783745
min   0.000000 0.000000 0.000000 0.002442
25%   1.000000 0.000000 0.000000 4.981579
50%   1.000000 1.000000 1.000000 10.069559
75%   1.000000 1.000000 1.000000 15.050018
max   1.000000 1.000000 1.000000 19.998709

      Previous Heart Problems Medication Use Stress Level \
count 8763.000000 8763.000000 8763.000000
mean  0.495835 0.498345 5.469702
std   0.500011 0.500026 2.859622
min   0.000000 0.000000 1.000000
25%   0.000000 0.000000 3.000000
50%   0.000000 0.000000 5.000000
75%   1.000000 1.000000 8.000000
max   1.000000 1.000000 10.000000

      Sedentary Hours Per Day Income BMI Triglycerides \
count 8763.000000 8763.000000 8763.000000 8763.000000
mean  5.993690 158263.181901 28.891446 417.677051
std   3.466359 80575.190806 6.319181 223.748137
min   0.001263 20062.000000 18.002337 30.000000
25%   2.998794 88310.000000 23.422985 225.500000
50%   5.933622 157866.000000 28.768999 417.000000
75%   9.019124 227749.000000 34.324594 612.000000
max   11.999313 299954.000000 39.997211 800.000000

      Physical Activity Days Per Week Sleep Hours Per Day Heart Attack Risk
count 8763.000000 8763.000000 8763.000000
mean  3.489672 7.023508 0.358211
std   2.282687 1.988473 0.479502
min   0.000000 4.000000 0.000000
25%   2.000000 5.000000 0.000000
50%   3.000000 7.000000 0.000000

```

75%	5.000000	9.000000	1.000000
max	7.000000	10.000000	1.000000

```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8763 entries, 0 to 8762
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Patient ID                           8763 non-null   object
1   Age                                   8763 non-null   int64
2   Sex                                   8763 non-null   object
3   Cholesterol                           8763 non-null   int64
4   Blood Pressure                        8763 non-null   object
5   Heart Rate                            8763 non-null   int64
6   Diabetes                              8763 non-null   int64
7   Family History                        8763 non-null   int64
8   Smoking                               8763 non-null   int64
9   Obesity                               8763 non-null   int64
10  Alcohol Consumption                   8763 non-null   int64
11  Exercise Hours Per Week               8763 non-null   float64
12  Diet                                  8763 non-null   object
13  Previous Heart Problems                8763 non-null   int64
14  Medication Use                         8763 non-null   int64
15  Stress Level                           8763 non-null   int64
16  Sedentary Hours Per Day                 8763 non-null   float64
17  Income                                 8763 non-null   int64
18  BMI                                    8763 non-null   float64
19  Triglycerides                          8763 non-null   int64
20  Physical Activity Days Per Week         8763 non-null   int64
21  Sleep Hours Per Day                    8763 non-null   int64
22  Country                                8763 non-null   object
23  Continent                              8763 non-null   object
24  Hemisphere                             8763 non-null   object
25  Heart Attack Risk                       8763 non-null   int64
dtypes: float64(3), int64(16), object(7)
memory usage: 1.7+ MB
```

2.0.1 It is first apparent here that the data doesnt contain any missing values

```
[8]: df["Heart Attack Risk"].value_counts()
```

```
[8]: Heart Attack Risk
0    5624
1    3139
Name: count, dtype: int64
```

3 Initial Cleaning

3.0.1 To get the maximum values out of my plot, I would like to first convert the existing categorical columns into numeric ones that can be visualized.

categorical_columns = "Patient ID", "Sex", "Blood Pressure", "Diet", "Country", "Continent,"Hemisphere"

```
[10]: df["Patient ID"].str[:3].value_counts() # Mostly unique, dropping the column
```

```
[10]: Patient ID
      MPX      6
      BGV      5
      OMJ      5
      RLR      5
      VAA      4
      ..
      AWA      1
      SUP      1
      MQX      1
      MVX      1
      XKA      1
      Name: count, Length: 6866, dtype: int64
```

```
[11]: df2 = df.drop(["Patient ID"],axis=1) #Patient ID
```

```
[12]: df.Sex.value_counts();
```

```
[13]: #Sex [One-Hot Encoding]
df2["Male"] = df.Sex.apply(lambda x: 1 if x == "Male" else 0)
df3 = df2.drop(["Sex"], axis=1)
```

```
[14]: # Blood Pressure
df3["Blood Pressure"]
split = pd.DataFrame(df3['Blood Pressure'].str.split('/').to_list(), columns =_
↳ ['Systolic', 'Diastolic'])
df4 = pd.concat([df3, split], axis=1)
df4.Systolic = df4.Systolic.astype(np.int32)
df4.Diastolic = df4.Diastolic.astype(np.int32)
df5 = df4.drop(['Blood Pressure'], axis = 1)
```

```
[15]: # Diet [Ordinal Encoding]
def encode_diet(x):
    if x == 'Healthy':
        return 0
    elif x == 'Average':
        return 1
    else:
```

```
return 2
```

```
df5['Bad_Diet'] = df5['Diet'].apply(lambda x: encode_diet(x))  
df6 = df5.drop(["Diet"], axis = 1)
```

One-Hot Encoding : “Country”, “Continent”, “Hemisphere”, would result in too many columns that contain information that is barely if at all to “Heart Attack Risk”. That is why I’m deciding to drop them.

```
[16]: df7 = df6.drop(["Country", "Continent", "Hemisphere"],axis = 1)
```

```
[17]: df7.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 8763 entries, 0 to 8762
```

```
Data columns (total 23 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	8763 non-null	int64
1	Cholesterol	8763 non-null	int64
2	Heart Rate	8763 non-null	int64
3	Diabetes	8763 non-null	int64
4	Family History	8763 non-null	int64
5	Smoking	8763 non-null	int64
6	Obesity	8763 non-null	int64
7	Alcohol Consumption	8763 non-null	int64
8	Exercise Hours Per Week	8763 non-null	float64
9	Previous Heart Problems	8763 non-null	int64
10	Medication Use	8763 non-null	int64
11	Stress Level	8763 non-null	int64
12	Sedentary Hours Per Day	8763 non-null	float64
13	Income	8763 non-null	int64
14	BMI	8763 non-null	float64
15	Triglycerides	8763 non-null	int64
16	Physical Activity Days Per Week	8763 non-null	int64
17	Sleep Hours Per Day	8763 non-null	int64
18	Heart Attack Risk	8763 non-null	int64
19	Male	8763 non-null	int64
20	Systolic	8763 non-null	int32
21	Diastolic	8763 non-null	int32
22	Bad_Diet	8763 non-null	int64

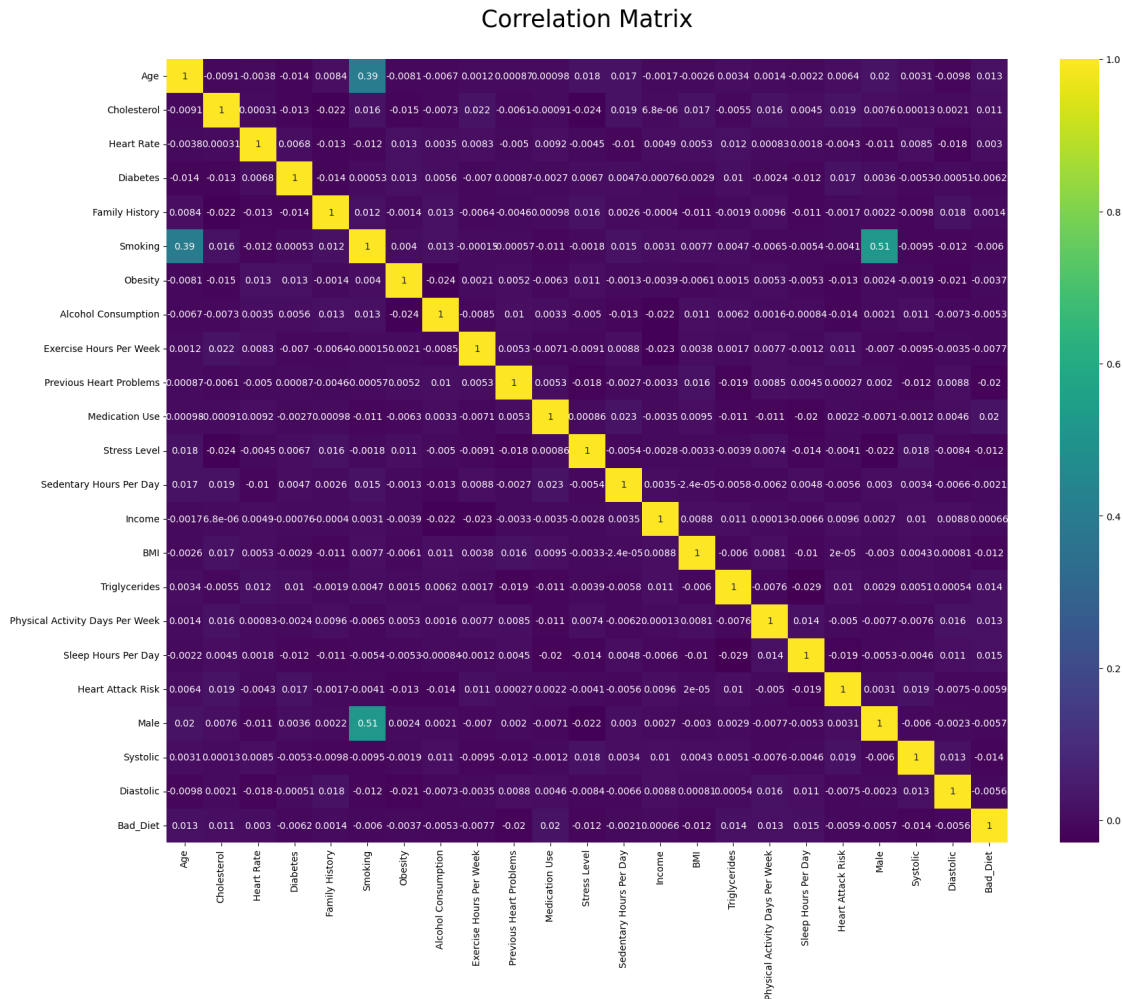
```
dtypes: float64(3), int32(2), int64(18)
```

```
memory usage: 1.5 MB
```

4 Exploratory Data Analysis

```
[18]: corr = df7.corr()
```

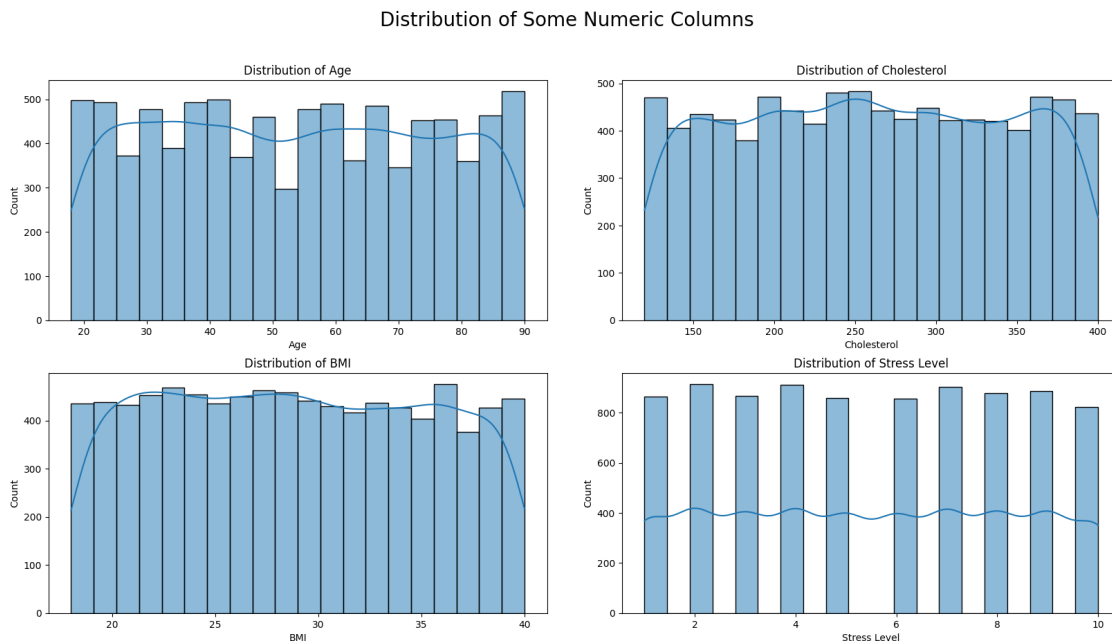
```
[19]: cmap = sns.color_palette("viridis", as_cmap=True)
fig,ax = plt.subplots(figsize=(20,15))
ax = sns.heatmap(corr,annot=True,cmap=cmap)
ax.set_title("Correlation Matrix\n",fontsize=25)
plt.show()
```



4.0.1 From the Heatmap we can see that there is only a clear relation between 2 columns:

Findings: - Smoking and Age, which makes sense & Smoking and Men. - Heart Attack risk has highest correlation with Diabetes, Cholesterol and Exercise Hours Per Week - Heart attack Risk is not much dependent on Sedentary Hours Per Day

```
[20]: # Distribution of Age
fig, ax = plt.subplots(2,2,figsize=(20,10))
fig.suptitle('Distribution of Some Numeric Columns',fontsize=20)
fig.subplots_adjust(hspace=0.22, wspace=0.15)
sns.histplot(ax=ax[0,0], data=df7["Age"],bins=20, kde=True).
    ↪set(title='Distribution of Age')
sns.histplot(ax=ax[0,1], data=df7["Cholesterol"],bins=20, kde=True).
    ↪set(title='Distribution of Cholesterol')
sns.histplot(ax=ax[1,0], data=df7["BMI"],bins=20, kde=True).
    ↪set(title='Distribution of BMI')
sns.histplot(ax=ax[1,1], data=df7["Stress Level"],bins=20, kde=True).
    ↪set(title='Distribution of Stress Level')
plt.show()
```



4.0.2 Now lets take a look at some of the divisions between the binary features

```
[21]: fig, ax = plt.subplots(2,4,figsize=(20,10))
fig.suptitle('Distribution of Some Categorical Columns',fontsize=20)
fig.subplots_adjust(hspace=0.22, wspace=0.2)
sns.countplot(ax=ax[0,0], data=df7,x = 'Diabetes').set(title='Diabetes')
sns.countplot(ax=ax[0,1], data=df7,x = 'Family History').set(title='Family_
    ↪History')
sns.countplot(ax=ax[0,2], data=df7,x = 'Smoking').set(title='Smoking')
sns.countplot(ax=ax[0,3], data=df7,x = 'Obesity').set(title='Obesity')
```

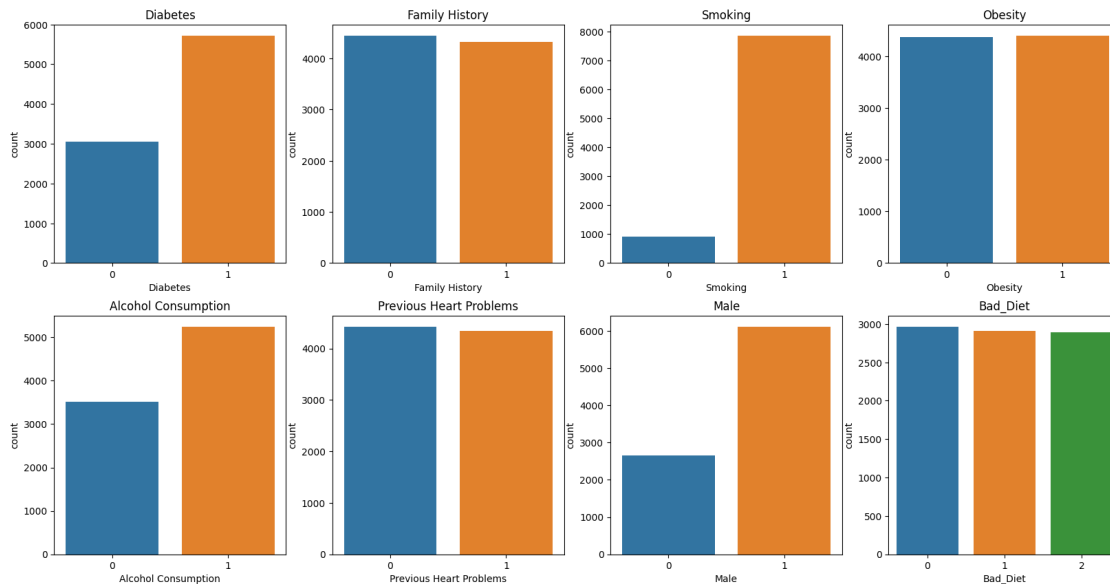


```

sns.countplot(ax=ax[1,0], data=df7,x ="Alcohol Consumption").set(title='Alcohol_
↳Consumption')
sns.countplot(ax=ax[1,1], data=df7,x ="Previous Heart Problems").
↳set(title='Previous Heart Problems')
sns.countplot(ax=ax[1,2], data=df7,x ="Male").set(title='Male')
sns.countplot(ax=ax[1,3], data=df7,x ="Bad_Diet").set(title='Bad_Diet')
plt.show()

```

Distribution of Some Categorical Columns



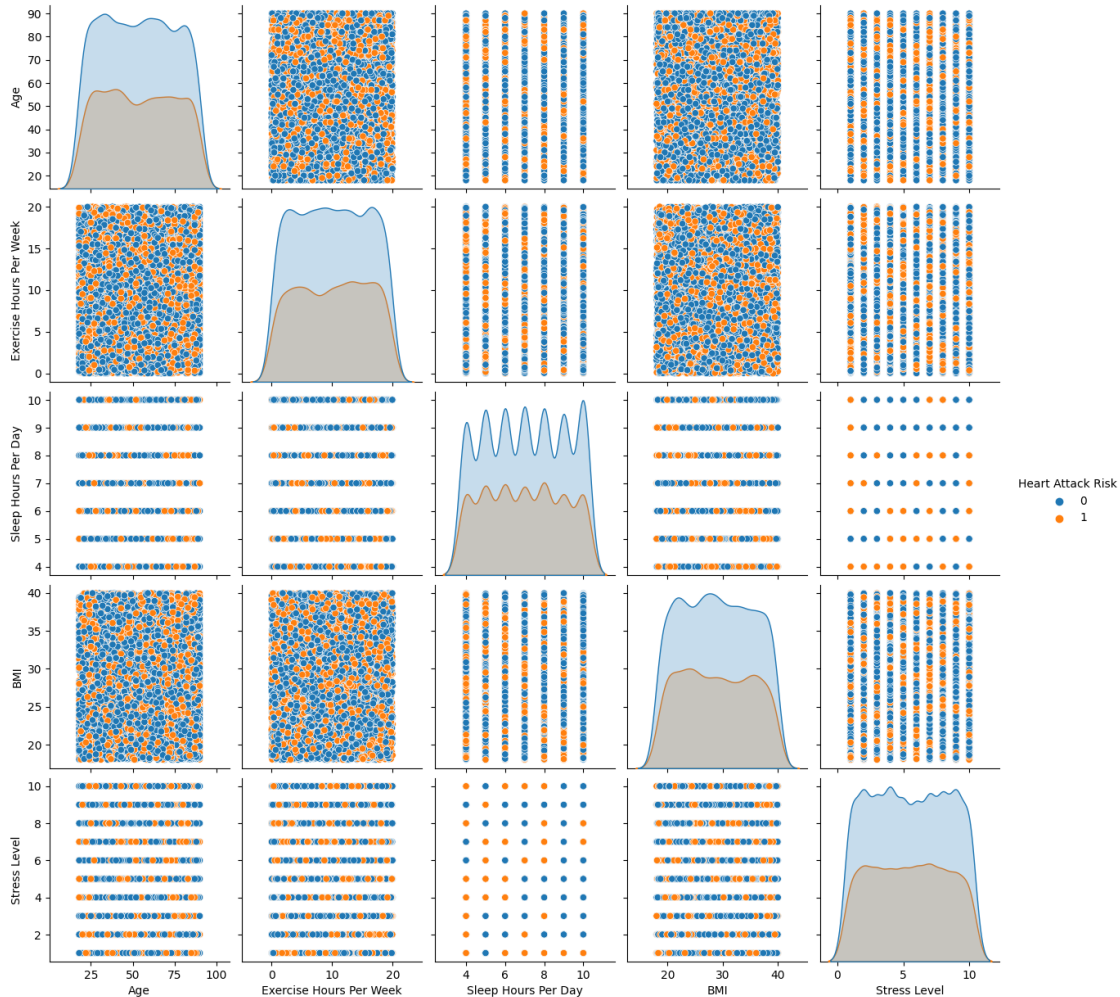
4.0.3 Since there is hardly any correlation in the data, scatter plots probably won't be very useful. However, the EDA that we can still do is look through some of the columns and compare the different labels

```

[22]: plt.figure(figsize=(10,10))
sns.pairplot(data=df7, vars=['Age', 'Exercise Hours Per Week', 'Sleep Hours Per_
↳Day', 'BMI', 'Stress Level'], hue='Heart Attack Risk', diag_kind='kde')
plt.show()

```

<Figure size 1000x1000 with 0 Axes>



4.0.4 We can see that the data is very uniformly distributed, whether or not the person is at risk of a heart disease

4.1 Detecting outliers

```
[23]: for i in df7.columns:
    Q1 = df7[i].quantile(0.25)
    Q3 = df7[i].quantile(0.75)
    IQR = Q3 - Q1
    threshold = 1.5
    outliers = df7[(df7[i] < Q1 - threshold * IQR) | (df7[i] > Q3 + threshold * IQR)]
    print(f"{i}<35>Outliers : {outliers.shape[0]}")
```

Age Outliers : 0
Cholesterol Outliers : 0

Heart Rate	Outliers : 0
Diabetes	Outliers : 0
Family History	Outliers : 0
Smoking	Outliers : 904
Obesity	Outliers : 0
Alcohol Consumption	Outliers : 0
Exercise Hours Per Week	Outliers : 0
Previous Heart Problems	Outliers : 0
Medication Use	Outliers : 0
Stress Level	Outliers : 0
Sedentary Hours Per Day	Outliers : 0
Income	Outliers : 0
BMI	Outliers : 0
Triglycerides	Outliers : 0
Physical Activity Days Per Week	Outliers : 0
Sleep Hours Per Day	Outliers : 0
Heart Attack Risk	Outliers : 0
Male	Outliers : 0
Systolic	Outliers : 0
Diastolic	Outliers : 0
Bad_Diet	Outliers : 0

```
[24]: df7.Smoking.value_counts() # Not really outliers
```

```
[24]: Smoking
1      7859
0       904
Name: count, dtype: int64
```

5 Data Preparation

```
[25]: X = df7.drop(["Heart Attack Risk"],axis = 1)
      y = df7["Heart Attack Risk"]
```

```
[26]: X.shape
```

```
[26]: (8763, 22)
```

```
[27]: y.shape
```

```
[27]: (8763,)
```

5.1 Scaling

Although not useful for every algorithm, it works for some and doesn't hurt the others. Since our data, is not heavy skewed on either side, just using sklearn's StandardScaler should be fine.

```
[28]: sc = StandardScaler()
X_scaled = sc.fit_transform(X)
X2 = pd.DataFrame(X_scaled, columns=X.columns, index = X.index)
```

```
[29]: X2.describe(); # Mean at 0, std = 1
```

5.2 Train test Split

```
[30]: X_train, X_test, y_train, y_test = train_test_split(X2, y, test_size=0.
↳ 3, shuffle=True, random_state=0)
```

```
[31]: print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(6134, 22)
```

```
(2629, 22)
```

```
(6134,)
```

```
(2629,)
```

6 Model Selection

```
[32]: models = {'LogisticRegression' : LogisticRegression(),
                'DecisionTreeClassifier' : DecisionTreeClassifier(),
                'RandomForestClassifier' : RandomForestClassifier(),
                'GradientBoostingClassifier' : GradientBoostingClassifier(),
                'CatBoostClassifier' : CatBoostClassifier(verbose=0)}
```

6.0.1 Baseline Models

```
[33]: def run_models(X_train, X_test, y_train, y_test):
    resultsdf = pd.DataFrame(columns = [
↳ ['Model', 'Train_Accuracy', 'Test_Accuracy', 'F1_Score', 'Roc_Auc'])
    for name, model in models.items():
        model.fit(X_train, y_train)
        print(f"{name} trained.")
        y_pred_train = model.predict(X_train)
        y_pred = model.predict(X_test)

        train_accuracy = accuracy_score(y_train, y_pred_train)
        test_accuracy = accuracy_score(y_test, y_pred)
        f1 = f1_score(y_test, y_pred)
        roc_auc = roc_auc_score(y_test, y_pred)
```

```

        results = {'Model': name, 'Train_Accuracy':train_accuracy,
        ↪ 'Test_Accuracy': test_accuracy, 'F1_Score': f1, 'Roc_Auc': roc_auc}
        resultsdf.loc[len(resultsdf)] = results
        return resultsdf

baseline = run_models(X_train,X_test,y_train,y_test)

```

LogisticRegression trained.
 DecisionTreeClassifier trained.
 RandomForestClassifier trained.
 GradientBoostingClassifier trained.
 CatBoostClassifier trained.

```
[34]: baseline.sort_values("F1_Score",ascending=False)
```

```
[34]:
```

	Model	Train_Accuracy	Test_Accuracy	F1_Score	\
1	DecisionTreeClassifier	1.000000	0.546215	0.382185	
4	CatBoostClassifier	0.828171	0.634081	0.099251	
2	RandomForestClassifier	1.000000	0.638266	0.066732	
3	GradientBoostingClassifier	0.664493	0.635983	0.030395	
0	LogisticRegression	0.643300	0.638266	0.000000	




	Roc_Auc
1	0.511944
4	0.508795
2	0.507745
3	0.501629
0	0.500000


Even though all these algorithms hav decent accuracy allround, an F1 score near 0 shoes that it really isn't predicting anything.

7 Feature Selection

```
[35]: for name, model in models.items():
        try:
            importances = model.feature_importances_
            indices = np.argsort(importances)[::-1]
            names = [X.columns[i] for i in indices]
            plt.figure(figsize=(15, 5))
            plt.title(f"{name} Feature Importance")
            plt.barh(range(X.shape[1]), importances[indices])
            plt.yticks(range(X.shape[1]), names)
            plt.gca().invert_yaxis()
            plt.show()
        except:
```

pass





7.0.1 Findings:

Most Important Features: - BMI - Exercise Hours Per Week - Triglycerides

Least Important Features: - Smoking - Diabetes - Previous Heart Problems - Family History

Lets take these features out and try again.

```
[36]: X_select = X2.drop(["Smoking","Diabetes","Previous Heart Problems", "Family_History", "Medication Use"],axis=1)
```

```
[37]: X_train_select,X_test_select,y_train_select,y_test_select = train_test_split(X_select,y,test_size=0.3,shuffle=True,random_state=0)
```

```
[38]: feature_selection = run_models(X_train_select,X_test_select,y_train,y_test)
```

LogisticRegression trained.

DecisionTreeClassifier trained.

RandomForestClassifier trained.

GradientBoostingClassifier trained.

CatBoostClassifier trained.

```
[39]: feature_selection.sort_values(by="F1_Score",ascending=False)
```

```
[39]:
```

	Model	Train_Accuracy	Test_Accuracy	F1_Score	\
1	DecisionTreeClassifier	1.000000	0.552301	0.387936	
4	CatBoostClassifier	0.829638	0.627235	0.095941	
2	RandomForestClassifier	1.000000	0.628376	0.059673	
3	GradientBoostingClassifier	0.664656	0.635223	0.026396	
0	LogisticRegression	0.643300	0.638266	0.000000	

	Roc_Auc
1	0.517623
4	0.503204
2	0.499314

```
3  0.500577
0  0.500000
```

Little Improvement

8 Removing Imbalance

```
[40]: y.value_counts()
```

```
[40]: Heart Attack Risk
0      5624
1      3139
Name: count, dtype: int64
```

There is a clear imbalance between 0s and 1s, in the target column making it so that a model can be ~63% accurate by just saying 0 everytime

```
[41]: 5624 - 3139
```

```
[41]: 2485
```

```
[42]: to_drop = y.where(lambda x: x==0).dropna().sample(2485,random_state = 99).index
      ↪#Not the best practice.
```

```
[43]: X_dropped = X_select.drop(to_drop,axis=0)
```

```
[44]: y_dropped = y.drop(to_drop)
```

```
[45]: X_train_dropped,X_test_dropped,y_train_dropped,y_test_dropped =
      ↪train_test_split(X_dropped,y_dropped,test_size=0.
      ↪3,shuffle=True,random_state=0)
```

```
[46]: dropped =
      ↪run_models(X_train_dropped,X_test_dropped,y_train_dropped,y_test_dropped)
```

LogisticRegression trained.
DecisionTreeClassifier trained.
RandomForestClassifier trained.
GradientBoostingClassifier trained.
CatBoostClassifier trained.

```
[47]: dropped
```

```
[47]:
```

	Model	Train_Accuracy	Test_Accuracy	F1_Score	\
0	LogisticRegression	0.521848	0.507431	0.493450	
1	DecisionTreeClassifier	1.000000	0.496815	0.495208	

2	RandomForestClassifier	1.000000	0.500531	0.486634
3	GradientBoostingClassifier	0.708921	0.501592	0.492707
4	CatBoostClassifier	0.945380	0.499469	0.505506

	Roc_Auc
0	0.507592
1	0.496846
2	0.500689
3	0.501699
4	0.499418

This helped the F1 Score come up for every model by a significant margin.

9 Correcting Overfitting [HyperParameter Tuning]

```
[48]: X_train_final = X_train_dropped
      X_test_final = X_test_dropped
      y_train_final = y_train_dropped
      y_test_final = y_test_dropped
```

From this point on I will just be focusing on 3 models. - Decision Tree Classifier - Random Forest Classifier - Categorical Boost Classifier

```
[49]: def test_model(model):
      resultsdf = pd.DataFrame(columns = [
      ↪ ['Train_Accuracy', "Test_Accuracy", "F1_Score", "Roc_Auc"])
      model.fit(X_train_final, y_train_final)
      y_pred_train = model.predict(X_train_final)
      y_pred = model.predict(X_test_final)
      train_accuracy = accuracy_score(y_train_final, y_pred_train)
      test_accuracy = accuracy_score(y_test_final, y_pred)
      f1 = f1_score(y_test_final, y_pred)
      roc_auc = roc_auc_score(y_test_final, y_pred)
      results = {'Train_Accuracy': train_accuracy, 'Test_Accuracy': test_accuracy,
      ↪ 'F1_Score': f1, 'Roc_Auc': roc_auc}
      resultsdf.loc[0] = results
      return resultsdf
```

9.0.1 Decision Tree Classifier

```
[50]: dt = DecisionTreeClassifier()
```

```
[51]: dt_results = test_model(dt)
      dt_results
```

```
[51]:   Train_Accuracy  Test_Accuracy  F1_Score   Roc_Auc
      0             1.0          0.501062  0.504219  0.501042
```

```
[52]: # dt_params = {"max_depth" : [3,5,7,9,11],
#               "criterion" : ["gini", "entropy", "log_loss"],
#               "min_samples_split" : [2,4,6,8]
#           }

# tuned_dt =
↳ GridSearchCV(dt,dt_params,verbose=0,cv=5,return_train_score=True,scoring='f1')
# tuned_dt.fit(X_train,y_train)
# tuned_dt.best_params_
```

```
[53]: best_dt = DecisionTreeClassifier(criterion='gini', max_depth = 11,
↳ min_samples_split =2) #Parameters from GridSearchCV above
dt_results = test_model(best_dt)
dt_results
```

```
[53]:   Train_Accuracy  Test_Accuracy  F1_Score   Roc_Auc
      0             0.63746          0.500531  0.355921  0.501737
```

Didn't really work

9.0.2 Random Forest Classifier

```
[54]: rf = RandomForestClassifier()
```

```
[55]: rf_results = test_model(rf)
rf_results
```

```
[55]:   Train_Accuracy  Test_Accuracy  F1_Score   Roc_Auc
      0             1.0          0.485669  0.479871  0.485742
```

```
[56]: # rf_params = {'n_estimators' : [10,20,30,50,100],
#               "max_depth" : [2,4,6,8]}

# tuned_rf =
↳ GridSearchCV(rf,rf_params,verbose=0,cv=5,return_train_score=False,scoring='f1')
# tuned_rf.fit(X_train,y_train)
# tuned_rf.best_params_
```

```
[57]: best_rf = RandomForestClassifier(max_depth = 8,n_estimators = 10) #Parameters
↳ from GridSearchCV above
rf_results = test_model(best_rf)
rf_results
```

```
[57]:   Train_Accuracy  Test_Accuracy  F1_Score   Roc_Auc
      0           0.742831        0.503185  0.497854  0.503255
```

Worked a bit. Probably resulted in better model.

9.0.3 Categorical Boost Classifier

```
[58]: cb = CatBoostClassifier(verbose=0)
```

```
[59]: cb_results = test_model(cb)
      cb_results
```

```
[59]:   Train_Accuracy  Test_Accuracy  F1_Score   Roc_Auc
      0           0.94538        0.499469  0.505506  0.499418
```

```
[60]: # cb_params={"depth" : [4,8,10],
#               'learning_rate' : [0.5,0.75],
#               'iterations' : [20,30,50],
#               "random_strength" : [2,5,10]
#           }

# tuned_cb =
↳ GridSearchCV(cb,cb_params,verbose=0,cv=5,return_train_score=False,scoring='f1')
# tuned_cb.fit(X_train,y_train)
# tuned_cb.best_params_
```

```
[61]: best_cb = CatBoostClassifier(depth = 8,iterations = 30,learning_rate=0.5,
↳ random_strength=5,verbose=0) #Parameters from GridSearchCV above
      cb_results = test_model(best_cb)
      cb_results
```

```
[61]:   Train_Accuracy  Test_Accuracy  F1_Score   Roc_Auc
      0           0.840237        0.518047  0.515475  0.518089
```

The best accuracy and f1 score of all the models.

Despite my best efforts to avoid overfitting, the models still tended to do it. When limiting max depth too much it came at a high cost of accuracy. Each of the Grid-SearchCVs done on the models aimed to maximize their f1_score. Even though it did improve it, the difference was very minute.

10 Conclusion

Given the data that I was, which mostly comprised of superficial attributes of a person, I couldn't predict the risk of a heart attack for a given person. Maybe if the features were more scientific, like the division of Cholesterol into HDL and LDL would have helped. Something that could have

improved the model more would be correctly balancing the labelled data, by not just removing a bunch of data. Furthermore, none of these methods used neural networks/deep learning models. This could be an interesting avenue to go down next time. As it stands, it is best to say that we shouldn't use this model to predict a person's risk of having a heart attack.